

Package ‘MST.PMDN’

June 14, 2025

Type Package

Title Deep Multivariate Skew t-Parsimonious Mixture Density Network

Version 0.1.0

Maintainer Alex J. Cannon <alex.cannon@ec.gc.ca>

Author Alex J. Cannon [aut, cre] (<<https://orcid.org/0000-0002-8025-3790>>)

Description Implements the Multivariate Skew t-Parsimonious Mixture Density Network (MST-PMDN), a distributional deep learning regression framework based on a mixture of MST distributions. Provides user-facing functions to define, train, predict, and sample from deep MST-PMDN models built with torch for R.

License GPL (>= 2)

Encoding UTF-8

Depends R (>= 3.5.0), torch (>= 0.13.0)

Imports coro, stats

LazyData true

LazyDataCompression xz

NeedsCompilation no

R topics documented:

MST.PMDN-package	2
define_mst_pmdn	4
loss_mst_pmdn	5
predict_mst_pmdn	6
sample_mst_pmdn	6
sample_mst_pmdn_df	7
train_mst_pmdn	8
t_cdf	11
t_cdf_fast	11
t_cdf_int	12
t_cdf_slow	12
wave_surge	13

Index	14
--------------	-----------

Description

The **MST.PMDN** package implements the deep Multivariate Skew t-Parsimonious Mixture Density Network (MST-PMDN), a distributional deep learning regression framework built on **torch** for R. The **MST.PMDN** framework represents complicated joint output distributions as mixtures of MST components. A volume (L)-shape (A)-orientation (D) (LAD) eigenvalue decomposition parameterization provides a tractable, interpretable, and parsimonious representation of the MST scale matrices, while explicit modeling of skewness and heavy tails can represent asymmetric behavior and tail dependence observed in real-world data (e.g., compound events and extremes). Parameters of a mixture of MST distributions that describe a multivariate output are estimated by training a deep learning model with two multi-modal input branches, one for tabular inputs and the other for (optional) image inputs. The two branches are provided as user-defined **torch** modules. Outputs from each are concatenated and passed through a dense fusion network, which then leads to the MST-PMDN head. In the absence of both branches, the tabular inputs are fed directly into the dense network.

Following the approach used in model-based clustering (**mclust**), scale matrices in the MST-PMDN head are represented using an LAD eigen-decomposition parameterization. LAD attributes, the nu (or degrees of freedom) parameter (n), and the alpha (or skewness) parameter (s) can be forced to be Variable or Equal between mixture components (plus Identity for A and D). For n and s, the model can also be constrained to emulate a multivariate Gaussian (or normal) (N) distribution. Different model types are specified by setting the argument `constraint = "EIINN", "VEVEV", etc.` where each letter position in the argument corresponds, respectively, to each of the LADns attributes. In the case of n, users can specify fixed (F) values for nu by passing an optional `fixed_nu` vector. If an element of `fixed_nu` is set to NA, then the value of nu for this component is learned by the network. Furthermore, values of mu (or means) (m), pi (or mixing coefficients) (x), volume-shape-orientation attributes (LAD), nu (n), and skewness (s) for the mixtures can be made to be independent of inputs by specifying any combination of `constant_attr = "m", "mx", ..., "LADmxns"`.

By combining appropriate values of `constraint` and `constant_attr`, MST-PMDN can emulate the Gaussian finite mixture models implemented by **mclust**, i.e., for unconditional density estimation or model-based clustering. Similarly, if the constraint on the nu parameter (n) is loosened (e.g., `constraint = "VVVEN"` with `constant_attr = "LADmxn"`), MST-PMDN can emulate model-based multivariate t clustering models provided by **teigen**. Going one step further, removing the constraint on the skewness parameter (s) (e.g., `constraint = "VVVEE"` with `constant_attr = "LADmxns"`) implements model-based MST clustering.

While it can be used for model-based density estimation and clustering tasks, the primary purpose of the **MST.PMDN** package is to implement likelihood-based deep generative models. With unconstrained or partially constrained `constant_attr`, the MST-PMDN framework allows parameters of the mixture of multivariate Gaussian, t, or skew t distributions to depend on tabular and image covariates via user-specified **torch** modules.

Details

Key user-facing functions include:

- `define_mst_pmdn`: construct a new MST-PMDN model object.
- `train_mst_pmdn`: fit an MST-PMDN model to training data.
- `predict_mst_pmdn`: generate predictive distributions or point predictions.
- `loss_mst_pmdn`: compute the negative log-likelihood loss for evaluation.
- `sample_mst_pmdn`: draw samples as `torch_tensor` objects.
- `sample_mst_pmdn_df`: draw samples returned as an R `data.frame`.

References

- Ambrogioni, L., Güçlü, U., van Gerven, M. A., & Maris, E. (2017). The kernel mixture network: A nonparametric method for conditional density estimation of continuous random variables. *arXiv:1705.07111*.
- Andrews, J. L., & McNicholas, P. D. (2012). Model-based clustering, classification, and discriminant analysis via mixtures of multivariate t-distributions: the t EIGEN family. *Statistics and Computing*, 22, 1021-1029.
- Azzalini, A., & Capitanio, A. (2003). Distributions generated by perturbation of symmetry with emphasis on a multivariate skew t-distribution. *Journal of the Royal Statistical Society Series B: Statistical Methodology*, 65(2), 367-389.
- Andrews, J. L., Wickins, J. R., Boers, N. M., & McNicholas, P. D. (2018). *teigen*: An R package for model-based clustering and classification via the multivariate t distribution. *Journal of Statistical Software*, 83, 1-32.
- Banfield, J. D., & Raftery, A. E. (1993). Model-based Gaussian and non-Gaussian clustering. *Biometrics*, 803-821.
- Celeux, G., & Govaert, G. (1995). Gaussian parsimonious clustering models. *Pattern Recognition*, 28(5), 781-793.
- Falbel D., & Luraschi, J. (2025). *torch*: Tensors and Neural Networks with 'GPU' Acceleration. R package version 0.14.2, <https://github.com/mlverse/torch>, <https://torch.mlverse.org/docs>.
- Fraley, C., & Raftery, A. E. (2002). Model-based clustering, discriminant analysis, and density estimation. *Journal of the American Statistical Association*, 97(458), 611-631.
- Fraley, C., & Raftery, A. E. (1998). How many clusters? Which clustering method? Answers via model-based cluster analysis. *The Computer Journal*, 41(8), 578-588.
- Lee, S., & McLachlan, G. J. (2014). Finite mixtures of multivariate skew t-distributions: some recent and new results. *Statistics and Computing*, 24, 181-202.
- Kingma, D. P., & Ba, J. (2015). Adam: a method for stochastic optimization. *Proceedings of the 3rd International Conference on Learning Representations, ICLR 2015, San Diego, CA, USA*. *arXiv:1412.6980*
- Klein, N. (2024). Distributional regression for data analysis. *Annual Review of Statistics and Its Application*, 11:321-346.
- Peel, D., & McLachlan, G.J. (2000). Robust mixture modelling using the t distribution. *Statistics and Computing* 10, 339–348.
- Srucca, L., Fop, M., Murphy, T. B., & Raftery, A. E. (2016). mclust 5: Clustering, classification and density estimation using Gaussian finite mixture models. *The R Journal*, 8(1), 289-317.
- Williams, P. M. (1996). Using neural networks to model conditional multivariate densities. *Neural Computation*, 8(4), 843-854.

define_mst_pmdn	<i>Construct an Untrained MST-PMDN Model</i>
-----------------	--

Description

Initializes a deep Multivariate Skew t-Parsimonious Mixture Density Network (MST-PMDN) as a **torch** module, ready for fitting via [train_mst_pmdn](#). The model predicts parameters of a mixture of multivariate skew t distributions using optional tabular and image feature extractors.

Usage

```
define_mst_pmdn(
    input_dim,
    output_dim,
    hidden_dim,
    n_mixtures,
    constraint      = "VVVNN",
    constant_attr   = "",
    activation      = nn_relu,
    drop_hidden    = 0,
    image_module    = NULL,
    tabular_module  = NULL,
    fixed_nu        = NULL,
    range_nu        = c(3, 50),
    max_alpha       = 5,
    min_vol_shape   = 1e-2,
    min_mix_weight  = 1e-4,
    jitter          = 1e-6
)
```

Arguments

input_dim	Integer. Number of features in the raw tabular input (used if tabular_module is NULL).
output_dim	Integer. Dimension d of the multivariate response.
hidden_dim	Integer vector. Sizes of hidden layers in the dense fusion network.
n_mixtures	Integer. Number of skew t mixture components M.
constraint	Character. Code for volume–shape–orientation, degrees of freedom nu, and skew constraints (e.g., “VVVFN”).
constant_attr	Character. Flags for parameters held constant (e.g., “m” for means, “s” for skew, etc.).
activation	Function. torch activation for hidden layers (default nn_relu).
drop_hidden	Numeric between 0 and 1. Dropout probability after each hidden layer.
image_module	A torch nn_module for image feature extraction, or NULL to omit.

tabular_module	A torch nn_module for tabular feature extraction, or NULL to use raw inputs.
fixed_nu	Numeric vector of length M, or NULL. If non-NULL, fixes degrees of freedom for each component.
range_nu	Numeric vector of length 2: c(min_nu, max_nu) to clamp learned nu.
max_alpha	Numeric. Absolute bound for skewness parameter.
min_vol_shape	Numeric. Minimum allowed value for volume (L) and shape (A) diagonal parameters.
min_mix_weight	Numeric. Floor for each mixture weight to avoid zero probabilities.
jitter	Numeric. Small ridge added to covariance diagonals for numerical stability.

Value

An untrained `mst_pmdn_model`, which is a **torch** nn_module encapsulating the MST-PMDN head. Use `train_mst_pmdn` to fit it to data.

loss_mst_pmdn	<i>Negative Log-Likelihood Loss for MST-PMDN</i>
---------------	--

Description

Computes the average negative log-likelihood under a mixture of multivariate skew t distributions for a batch of examples. The function extracts mixture weights, component means, Cholesky scale factors, degrees of freedom, and skewness parameters from output, then evaluates the MST-PMDN log-density at target, finally averaging (and negating) to produce the loss. The `nu_switch` argument determines at what nu value the code uses the faster (but approximate) versus slower CDF routines.

Usage

```
loss_mst_pmdn(output, target, nu_switch = 20)
```

Arguments

output	A list as returned by a forward pass of an <code>mst_pmdn_model</code> , containing at least the following named <code>torch_tensor</code> elements: <code>pi</code> , <code>mu</code> , <code>scale_chol</code> , <code>nu</code> , and <code>alpha</code> .
target	A <code>torch_tensor</code> or numeric matrix of true responses, with shape <code>[batch_size, d]</code> .
nu_switch	Numeric scalar. Degrees-of-freedom threshold for switching between the <code>t_cdf_fast</code> and <code>t_cdf_slow</code> implementations of the Student t CDF.

Value

A single-element `torch_tensor` containing the batch-averaged negative log-likelihood.

predict_mst_pmdn	<i>Generate predictions from a trained MST-PMDN model</i>
------------------	---

Description

Wraps a trained `mst_pmdn_model` in evaluation mode, converts inputs to `torch_tensor` if necessary, and runs a forward pass without tracking gradients to produce mixture-distribution parameters.

Usage

```
predict_mst_pmdn(model, new_inputs, image_inputs = NULL, device = "cpu")
```

Arguments

<code>model</code>	An <code>mst_pmdn_model</code> object (a torch module), typically returned by <code>train_mst_pmdn</code> .
<code>new_inputs</code>	A numeric matrix or <code>torch_tensor</code> of tabular predictors. If not already a tensor, it is converted via <code>torch_tensor</code> on the specified device.
<code>image_inputs</code>	Optional 4D array or <code>torch_tensor</code> of image data for models with an <code>image_module</code> . If provided, must match the module's expected shape; otherwise set to <code>NULL</code> .
<code>device</code>	Character; the torch device to use (e.g. "cpu" or "cuda"). Defaults to "cpu".

Value

A named list containing the output of the model's forward pass:

- `pi` Mixture weights tensor of shape `[batch_size, M]`.
- `mu` Component means tensor of shape `[batch_size, M, d]`.
- `scale_chol` Cholesky scale factors tensor of shape `[batch_size, M, d, d]`.
- `nu` Degrees-of-freedom tensor of shape `[batch_size, M]`.
- `alpha` Skewness parameters tensor of shape `[batch_size, M, d]`.
- `L, A, D` Volume, shape, and orientation decomposition tensors for each component.

sample_mst_pmdn	<i>Draw raw tensor samples from an MST-PMDN predictive mixture</i>
-----------------	--

Description

Performs mixture sampling from the multivariate skew t predictive distribution encapsulated in `mdn_output`. For each of the `num_samples` draws, a mixture component is sampled according to the mixture weights `pi`, and then a skew- t variate is generated via a Gaussian plus Gamma construction. The output tensors are permuted so that the first dimension indexes the draw number.

Usage

```
sample_mst_pmdn(mdn_output, num_samples = 1, device = "cpu")
```

Arguments

mdn_output A list as returned by `predict_mst_pmdn` or a forward pass of an `mst_pmdn_model`, containing at minimum the following named `torch_tensor` elements: `pi`, `mu`, `scale_chol`, `nu`, and `alpha`.

num_samples Integer; number of random draws to generate per input case. Defaults to 1.

device Character; a **torch** device on which to perform sampling (e.g. "cpu" or "cuda"). Defaults to "cpu".

Value

A list with components:

samples A `torch_tensor` of shape `c(num_samples, batch_size, d)`, giving the sampled responses.

components A `torch_tensor` of shape `c(num_samples, batch_size)`, containing the index (1–M) of the mixture component used for each draw.

sample_mst_pmdn_df	<i>Draw MST-PMDN predictive samples as a data frame</i>
--------------------	---

Description

For each of `num_samples` draws per case, a mixture component is sampled according to the weights `pi`, and then a multivariate skew-*t* variate is generated via the usual Gaussian + Gamma construction. The results are returned in long format as an R `data.frame` with one row per draw, including the sampled values for each response dimension, the input case index, the draw index, and the component chosen.

Usage

```
sample_mst_pmdn_df(mdn_output, num_samples = 1, device = "cpu")
```

Arguments

mdn_output A list as returned by `predict_mst_pmdn` or a forward pass of an `mst_pmdn_model`, containing at minimum the following named `torch_tensor` elements: `pi`, `mu`, `scale_chol`, `nu`, and `alpha`.

num_samples Integer; number of random draws to generate per input case. Defaults to 1.

device Character; the torch device on which sampling is performed (e.g. "cpu" or "cuda"). Defaults to "cpu".

Value

A `data.frame` with `num_samples * B` rows (where `B` = number of input cases) and `d + 3` columns:

`V1, ..., Vd` Numeric; sampled values for each of the d response dimensions.

`row` Integer; index of the input case (1 through `B`).

`draw` Integer; draw number (1 through `num_samples`).

`comp` Factor; mixture component index (1 through `M`) used for that draw.

train_mst_pmdn	<i>Train a Deep Multivariate Skew t Parsimonious Mixture Density Network</i>
----------------	---

Description

Fits an MST-PMDN model by minimizing the negative log-likelihood on the supplied data, with support for optional image inputs, checkpointing, early stopping, gradient clipping, weight decay, dropout, and learning-rate scheduling.

Usage

```
train_mst_pmdn(
  inputs,
  outputs,
  hidden_dim,
  n_mixtures,
  constraint      = "VVVNN",
  constant_attr   = "",
  fixed_nu        = NULL,
  range_nu        = c(3, 50),
  nu_switch       = 20,
  max_alpha       = 5,
  min_vol_shape   = 1e-2,
  min_mix_weight  = 1e-4,
  jitter          = 1e-6,
  activation       = nn_tanh,
  epochs          = 500,
  lr              = 0.001,
  batch_size      = 16,
  max_norm        = 1,
  drop_hidden     = 0,
  wd_image        = 0,
  wd_tabular      = 0,
  checkpoint_interval = 10,
  checkpoint_path  = "checkpoint.pt",
  resume_from_checkpoint = FALSE,
  model           = NULL,
```



```

    early_stopping_patience= 50,
    validation_split          = 0.2,
    custom_split              = NULL,
    scheduler_step            = 50,
    scheduler_gamma           = 0.5,
    image_inputs              = NULL,
    image_module              = NULL,
    tabular_module            = NULL,
    device                    = "cpu"
)

```

Arguments

inputs	Numeric matrix or torch_tensor of predictors (rows = cases).
outputs	Numeric matrix or torch_tensor of responses (rows = cases).
hidden_dim	Integer vector giving the sizes of the hidden layers in the fusion network.
n_mixtures	Integer; number of mixture components M.
constraint	Character; MST-PMDN constraint code (volume–shape–orientation, degrees of freedom nu, and skew), e.g. "VVVFN".
constant_attr	Character; flags for parameters held constant ("m" for means, "x" for mixing coefficients, "L" for volume, "A" for shape, "D" for orientation, "n" for degrees of freedom nu, and "s" for skew).
fixed_nu	Numeric vector of length M, or NULL; if non-NULL, fixes degrees of freedom nu per component.
range_nu	Numeric vector of length 2: clamp range for nu.
nu_switch	Numeric; nu threshold for switching between fast vs slow CDF routines.
max_alpha	Numeric; maximum absolute skewness parameter.
min_vol_shape	Numeric; minimum allowed volume (L) and shape (A) diagonal values.
min_mix_weight	Numeric; minimum mixture weight per component.
jitter	Numeric; small ridge added to covariance diagonals for stability.
activation	Activation function for hidden layers (e.g., nn_tanh, nn_relu).
epochs	Integer; maximum number of training epochs.
lr	Numeric; learning rate for the optimizer.
batch_size	Integer; mini-batch size.
max_norm	Numeric; maximum gradient norm for clipping.
drop_hidden	Numeric in [0,1]; dropout probability after each hidden layer.
wd_image	Numeric; weight decay (L2 penalty) for image_module parameters.
wd_tabular	Numeric; weight decay for tabular_module parameters.
checkpoint_interval	Integer; number of epochs between saved checkpoints.
checkpoint_path	Character; file path for saving model checkpoints.

resume_from_checkpoint	Logical; if TRUE, loads and resumes from an existing checkpoint.
model	An existing mst_pmdn_model to continue training, or NULL to initialize a new one.
early_stopping_patience	Integer; epochs with no validation improvement before stopping.
validation_split	Numeric in [0,1]; fraction of cases held out for validation if custom_split is NULL.
custom_split	Index or logical vector, or list with train and validation elements; overrides validation_split.
scheduler_step	Integer; epochs between learning-rate decay steps.
scheduler_gamma	Numeric; multiplicative factor for learning-rate decay.
image_inputs	Optional array or torch_tensor of image data (4D) for multimodal models.
image_module	Optional nn_module for processing image_inputs.
tabular_module	Optional nn_module for processing inputs.
device	Character; torch device to use (e.g., "cpu" or "cuda").

Details

Splits the data into training/validation sets (unless custom_split is provided), initializes or loads the model, and then runs an Adam optimization loop with optional gradient clipping, weight decay, dropout, early stopping, checkpointing, and learning-rate scheduling. After training, the best-performing model (by validation or training loss) is reloaded before returning.

Value

A list with components:

- model** The trained mst_pmdn_model (torch module).
- train_loss_history** Numeric vector of training losses per epoch.
- val_loss_history** Numeric vector of validation losses per epoch, or NULL if none.
- best_train_epoch** Epoch number with lowest training loss (if no validation split).
- best_train_loss** Lowest training loss achieved (if no validation split).
- best_val_epoch** Epoch number with lowest validation loss (if used).
- best_val_loss** Lowest validation loss achieved (if used).
- final_epoch** Last completed epoch.
- train_indices** Indices of cases used for training.
- val_indices** Indices of cases used for validation.

t_cdf

Unified interface to Student t CDF

Description

Conditionally dispatches to `t_cdf_fast` or `t_cdf_slow`, providing a single API for computing the Student t cumulative distribution in a manner compatible with torch computation graphs. Note: will be deprecated once `torch.distributions.studentT` has been ported to torch for R.

Usage

```
t_cdf(z, nu, nu_switch = 20)
```

Arguments

z	Numeric vector or <code>torch_tensor</code> of quantiles.
nu	Degrees of freedom (scalar or tensor).
nu_switch	Value of nu below/above which <code>t_cdf_slow</code> / <code>t_cdf_fast</code> are used.

Value

A `torch_tensor` of CDF values.

t_cdf_fast

Fast approximation of the Student t CDF

Description

Computes a differentiable “fast” approximation to the cumulative distribution function of the Student t distribution, suitable for use in **torch** computation graphs.

Usage

```
t_cdf_fast(z, nu)
```

Arguments

z	Numeric vector or <code>torch_tensor</code> of quantiles.
nu	Degrees of freedom (scalar or tensor).

Value

A `torch_tensor` of the same shape as z, containing CDF values.

t_cdf_int

Numerical integration-based Student t CDF

Description

Computes the Student t cumulative distribution function by numerically integrating the probability density function.

Usage

```
t_cdf_int(t_val, nu, num_integration_points = 1000L)
```

Arguments

t_val Numeric vector or [torch_tensor](#) of quantiles at which to evaluate the CDF.
 nu Degrees of freedom (scalar or tensor).
 num_integration_points Integer; number of points to use in the numerical integration. Defaults to 1000L.

Value

A [torch_tensor](#) of CDF values corresponding to t_val.

t_cdf_slow

Accurate but slower Student t CDF approximation

Description

Computes a more accurate but computationally slower approximation to the Student t CDF using pt and finite differences for the gradient.

Usage

```
t_cdf_slow(z, nu)
```

Arguments

z Numeric vector or [torch_tensor](#) of quantiles.
 nu Degrees of freedom (scalar or tensor).

Value

A [torch_tensor](#) of CDF values corresponding to z.

wave_surge

*Wave–Surge Daily-Maximum Subset***Description**

A subset of the CCCRIS node 181947 daily maximum wave and surge data and covariates.

Usage

```
data(wave_surge)
```

Format

A named list with three components:

x Numeric array of tabular predictors (harmonics of the seasonal cycle and lag-1 wave/surge data).

x_image Numeric array of image predictors (sea level pressure and sea level pressure gradient).

y Numeric matrix of scaled responses (daily maximum significant wave height and surge).

Source

CCCRIS node 181947 (cccris.ca) and ERA5 reanalysis

Index

- * **datasets**

- wave_surge, [13](#)

- * **package**

- MST.PMDN-package, [2](#)

data.frame, [3](#), [7](#), [8](#)

define_mst_pmdn, [3](#), [4](#)

loss_mst_pmdn, [3](#), [5](#)

MST.PMDN-package, [2](#)

predict_mst_pmdn, [3](#), [6](#), [7](#)

sample_mst_pmdn, [3](#), [6](#)

sample_mst_pmdn_df, [3](#), [7](#)

t_cdf, [11](#)

t_cdf_fast, [5](#), [11](#), [11](#)

t_cdf_int, [12](#)

t_cdf_slow, [5](#), [11](#), [12](#)

torch_tensor, [3](#), [5–7](#), [9](#), [11](#), [12](#)

train_mst_pmdn, [3–6](#), [8](#)

wave_surge, [13](#)